

class07 Clustering BIMM143 Scott MacLeod

Scott MacLeod

Clustering Methods

The broad goal here is to find groupings (clusters) in your input data.

Kmeans

First, let's make up some data to cluster.


We are going to use `rnorm()` in order to make up some numbers. For example:

```
rnorm(5)
```

```
[1] 1.287272 1.759522 2.963176 2.253911 1.270339
```

But, we are going to use 1,000 and turn it into a histogram!

```
x <- rnorm(1000)
hist(x)
```



Make a vector of length 60 with 30 points centered at -3 and 30 points centered at +3

```
tmp <- c(rnorm(30, mean=-3), rnorm(30, mean = 3))
tmp
```

```
[1] -2.3712600 -2.3881590 -2.4632339 -1.3430033 -2.0695799 -3.2314319
[7] -1.9028718 -2.3344299 -3.9799338 -2.6991218 -2.7234767 -3.1024425
[13] -3.0829499 -3.5722274 -4.6456327 -3.9673174 -2.4413441 -4.7840010
[19] -4.7045100 -2.9898411 -2.3348882 -4.6237639 -2.9815985 -4.2668254
[25] -3.6732881 -2.7657798 -1.4943215 -1.1313594 -4.3114108 -2.9491492
[31]  3.2385677  0.6859490  3.1219878  1.6210932  2.9014646  3.5444786
[37]  1.5461177  2.4624231  1.8820063  1.9835510  2.8747541  3.7701765
[43]  2.4487741  3.2755139  3.1310713  3.1805380  2.8372498  3.5611826
[49]  0.3714007  4.1804710  1.4335416  3.9345738  4.3826019  2.2796438
[55]  3.3306641  4.4146861  3.0408435  4.3314404  1.4983177  3.1346905
```


I will now make a small x and y dataset with 2 groups of points. Basically going to take the reverse of it! We are going to use the `rev()` function.

```
x <- cbind(x=tmp, y=rev(tmp))
x
```

	x	y
[1,]	-2.3712600	3.1346905
[2,]	-2.3881590	1.4983177
[3,]	-2.4632339	4.3314404
[4,]	-1.3430033	3.0408435
[5,]	-2.0695799	4.4146861
[6,]	-3.2314319	3.3306641
[7,]	-1.9028718	2.2796438
[8,]	-2.3344299	4.3826019
[9,]	-3.9799338	3.9345738
[10,]	-2.6991218	1.4335416
[11,]	-2.7234767	4.1804710
[12,]	-3.1024425	0.3714007
[13,]	-3.0829499	3.5611826
[14,]	-3.5722274	2.8372498
[15,]	-4.6456327	3.1805380
[16,]	-3.9673174	3.1310713
[17,]	-2.4413441	3.2755139
[18,]	-4.7840010	2.4487741
[19,]	-4.7045100	3.7701765
[20,]	-2.9898411	2.8747541
[21,]	-2.3348882	1.9835510
[22,]	-4.6237639	1.8820063
[23,]	-2.9815985	2.4624231
[24,]	-4.2668254	1.5461177
[25,]	-3.6732881	3.5444786
[26,]	-2.7657798	2.9014646
[27,]	-1.4943215	1.6210932
[28,]	-1.1313594	3.1219878
[29,]	-4.3114108	0.6859490
[30,]	-2.9491492	3.2385677
[31,]	3.2385677	-2.9491492
[32,]	0.6859490	-4.3114108
[33,]	3.1219878	-1.1313594
[34,]	1.6210932	-1.4943215
[35,]	2.9014646	-2.7657798
[36,]	3.5444786	-3.6732881
[37,]	1.5461177	-4.2668254
[38,]	2.4624231	-2.9815985
[39,]	1.8820063	-4.6237639
[40,]	1.9835510	-2.3348882
[41,]	2.8747541	-2.9898411
[42,]	3.7701765	-4.7045100

```
[43,] 2.4487741 -4.7840010
[44,] 3.2755139 -2.4413441
[45,] 3.1310713 -3.9673174
[46,] 3.1805380 -4.6456327
[47,] 2.8372498 -3.5722274
[48,] 3.5611826 -3.0829499
[49,] 0.3714007 -3.1024425
[50,] 4.1804710 -2.7234767
[51,] 1.4335416 -2.6991218
[52,] 3.9345738 -3.9799338
[53,] 4.3826019 -2.3344299
[54,] 2.2796438 -1.9028718
[55,] 3.3306641 -3.2314319
[56,] 4.4146861 -2.0695799
[57,] 3.0408435 -1.3430033
[58,] 4.3314404 -2.4632339
[59,] 1.4983177 -2.3881590
[60,] 3.1346905 -2.3712600
```

```
plot(x)
```



We are going to run `kmeans()`

```
k <- kmeans(x, centers = 2)
k
```

K-means clustering with 2 clusters of sizes 30, 30

Cluster means:

```
      x      y
1  2.813326 -3.044305
2 -3.044305  2.813326
```

Clustering vector:

```
[1] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 1 1 1 1 1 1 1 1
[39] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
```

Within cluster sum of squares by cluster:

```
[1] 63.37361 63.37361
(between_SS / total_SS =  89.0 %)
```

Available components:

```
[1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss"
[6] "betweenss"    "size"         "iter"         "ifault"
```

Q. From your result object k how many points are in each cluster?

```
k$size
```

```
[1] 30 30
```

Q. What “component” of your results object details the cluster membership?

```
k$cluster
```

```
[1] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 1 1 1 1 1 1 1 1
[39] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
```


Q. Cluster centers?

```
k$centers
```

	x	y
1	2.813326	-3.044305
2	-3.044305	2.813326

Q. Plot of our clustering results?

```
plot(x, col=k$cluster)
points(k$centers, col=5, pch=15, cex=2)
```



We can also cluster into 4 groups!

```
#kmeans
k4 <- kmeans(x, centers = 4)
k4
```

K-means clustering with 4 clusters of sizes 5, 18, 30, 7

Cluster means:

	x	y
1	-4.217689	1.386850
2	-3.163721	3.471475

```
3  2.813326 -3.044305
4 -1.899104  2.139854
```

Clustering vector:

```
[1] 2 4 2 4 2 2 4 2 2 4 2 1 2 2 2 2 2 1 2 2 4 1 2 1 2 2 4 4 1 2 3 3 3 3 3 3 3
[39] 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3
```


Within cluster sum of squares by cluster:

```
[1] 4.661207 16.113605 63.373607 5.131554
(between_SS / total_SS = 92.3 %)
```

Available components:

```
[1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss"
[6] "betweenss"    "size"         "iter"         "ifault"
```

```
#plot results
plot(x, col=k4$cluster)
```



A big limitation of `kmeans()` is that it does wht you ask even if you ask for silly clusters!

Hierarchical Clustering

The main base R function for Hierarchical Clustering is `hclust()`. Unlike `kmeans()` you can not just pass your data as input. You first need to calculate a distance matrix.

```
d <- dist(x)
hc <- hclust(d)
hc
```


Call:

```
hclust(d = d)
```

```
Cluster method   : complete
Distance          : euclidean
Number of objects: 60
```

Use `plot()` to view results.


```
plot(hc)
abline(h=8, col=5)
```




```
grps <- cutree(hc, h=8)
grps
```

[1] 1 2 2 2 2 2 2 2 2 2
[39] 2

```
plot(x, col=grps)
```



```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url, row.names=1)
x
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139
Fresh_potatoes	720	874	566	1033
Fresh_Veg	253	265	171	143
Other_Veg	488	570	418	355
Processed_potatoes	198	203	220	187
Processed_Veg	360	365	337	334
Fresh_fruit	1102	1137	957	674
Cereals	1472	1582	1462	1494
Beverages	57	73	53	47
Soft_drinks	1374	1256	1572	1506
Alcoholic_drinks	375	475	458	135
Confectionery	54	64	62	41

Q1. How many rows and columns are in your new data frame named `x`? What R functions could you use to answer this questions?

```
nrow(x)
```

```
[1] 17
```

```
ncol(x)
```

```
[1] 4
```


There are 17 rows and 4 columns. You can also look at the “environment” tab in the top right of the window.

Q2. Which approach to solving the ‘row-names problem’ mentioned above do you prefer and why? Is one approach more robust than another under certain circumstances?


I like using the `row.names()` function versus subtracting the rows. If you kept subtracting the columns, eventually you would run out of columns if you kept running the code. The `row.names()` approach is more robust and leaves less room for accidents in the future.

Q3. Changing what optional argument in the above `barplot()` function results in the following plot?

```
barplot(as.matrix(x), beside=T, col=rainbow(nrow(x)))
```




```
barplot(as.matrix(x), beside=F, col=rainbow(nrow(x)))
```



By changing the “Beside” from True to False, you can change the plots. If it is false, the columns of heights are portrayed as stacked bars, and if true the columns are portrayed as juxtaposed bars.

Q5. Generating all pairwise plots may help somewhat. Can you make sense of the following code and resulting figure? What does it mean if a given point lies on the diagonal for a given plot?

```
pairs(x, col=rainbow(10), pch=16)
```



This would be useful for a small set of Data. This Paris Plot is comparing all of the 4 countries. If we wanted to look further, we could color code the dots and see which foods/bevs are being compared.

##PCA to the rescue

The main “base” R function for PCA is called `prcomp()`. We are going to use `t()` to transpose the data set. Then we took the pcr and then made a summary table of the results.

```
pca <- prcomp(t(x))
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4
Standard deviation	324.1502	212.7478	73.87622	4.189e-14
Proportion of Variance	0.6744	0.2905	0.03503	0.000e+00
Cumulative Proportion	0.6744	0.9650	1.00000	1.000e+00

Q How much variance is captured in two PCs?

96.5% is captured in two PCs. Look at the cumulative tab in the table above :)

Now to make our main “PC score plot” (a.k.a “PC1 vs. PC2”, or “PC plot” or “Ordination plot”). This thing has a lot of different names.

```
attributes(pca)
```

```
$names
```

```
[1] "sdev"      "rotation" "center"    "scale"     "x"
```

```
$class
```

```
[1] "prcomp"
```


We are after the `pca$x` result component to make our main PCA plot.

```
pca$x
```

	PC1	PC2	PC3	PC4
England	-144.99315	2.532999	-105.768945	2.842865e-14
Wales	-240.52915	224.646925	56.475555	7.804382e-13
Scotland	-91.86934	-286.081786	44.415495	-9.614462e-13
N.Ireland	477.39164	58.901862	4.877895	1.448078e-13

```
mycols <- c("orange","red","blue","darkgreen")
```

```
plot(pca$x[,1], pca$x[,2], col=mycols, pch=16, xlab="PC1 (67.4%)", ylab="PC2(29%)")
```



Another important result from PCA is how the original variables (in this case: the foods) contributed to the PCs.

This is contained in the `pca$rotation()` object - people often call this the “loadings” or “contributions” to the PCs.

```
pca$rotation
```

	PC1	PC2	PC3	PC4
Cheese	-0.056955380	-0.016012850	-0.02394295	-0.691718038
Carcass_meat	0.047927628	-0.013915823	-0.06367111	0.635384915
Other_meat	-0.258916658	0.015331138	0.55384854	0.198175921
Fish	-0.084414983	0.050754947	-0.03906481	-0.015824630
Fats_and_oils	-0.005193623	0.095388656	0.12522257	0.052347444
Sugars	-0.037620983	0.043021699	0.03605745	0.014481347
Fresh_potatoes	0.401402060	0.715017078	0.20668248	-0.151706089
Fresh_Veg	-0.151849942	0.144900268	-0.21382237	0.056182433
Other_Veg	-0.243593729	0.225450923	0.05332841	-0.080722623
Processed_potatoes	-0.026886233	-0.042850761	0.07364902	-0.022618707
Processed_Veg	-0.036488269	0.045451802	-0.05289191	0.009235001
Fresh_fruit	-0.632640898	0.177740743	-0.40012865	-0.021899087
Cereals	-0.047702858	0.212599678	0.35884921	0.084667257
Beverages	-0.026187756	0.030560542	0.04135860	-0.011880823
Soft_drinks	0.232244140	-0.555124311	0.16942648	-0.144367046
Alcoholic_drinks	-0.463968168	-0.113536523	0.49858320	-0.115797605
Confectionery	-0.029650201	-0.005949921	0.05232164	-0.003695024

We can make a plot along PC1. This one isn't as pretty as it could be, but we have the general idea.

```
library(ggplot2)

contrib <- as.data.frame(pca$rotation)
ggplot(contrib) +
  aes(PC1, rownames(contrib)) +
  geom_col()
```

